
multiverse

multiverse

Natural Language Processing

Data Science Unit 4

Before we start...

- ➔ Make sure you are comfortable
- ➔ Have water and maybe a strong coffee handy
- ➔ If you need a break… take it!
- ➔ If you need a stretch – please go ahead!
- ➔ Please mute yourselves if you are not talking
- ➔ Have your video on at all times

...and let's get started!

In this session we will...

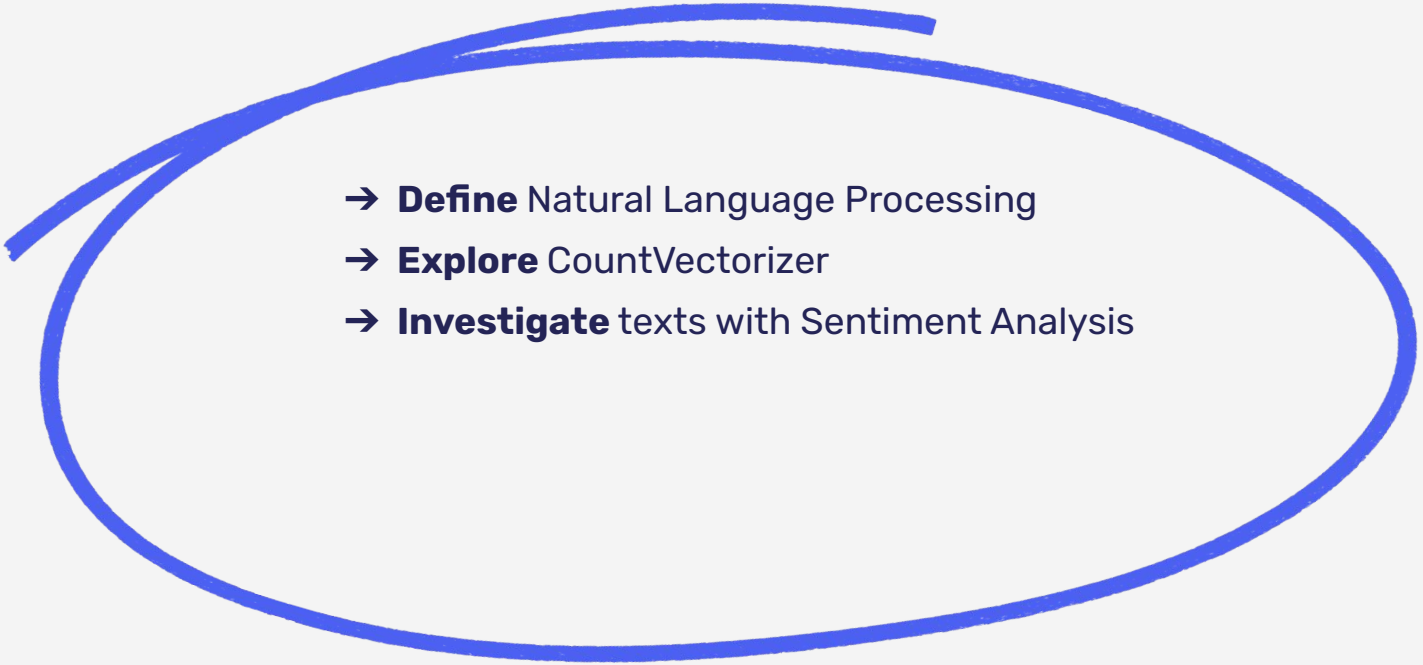

- ➔ **Define** Natural Language Processing
- ➔ **Explore** CountVectorizer
- ➔ **Investigate** texts with Sentiment Analysis

multiverse

NLP

NLP

Analysis

NLP techniques provide tools to allow us to understand and analyse large amounts of text. For example:

- ➔ Analyse the positivity/negativity of comments on different websites.
- ➔ Extract keywords from meeting notes and visualise how meeting topics change over time.

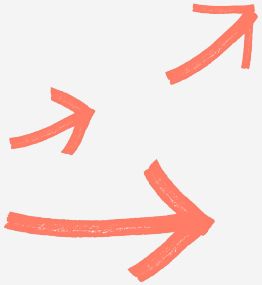

Vectorising for machine learning

When building a machine learning model, we typically must transform our data into numeric features. This process of transforming non-numeric data such as natural language into numeric features is called vectorisation. For example:

- ➔ Understanding related words. Using stemming, NLP lets us know that "swim", "swims", and "swimming" all refer to the same base word. This allows us to reduce the number of features used in our model.
- ➔ Identifying important and unique words. Using TF-IDF (term frequency-inverse document frequency), we can identify which words are most likely to be meaningful in a document.

What is it?

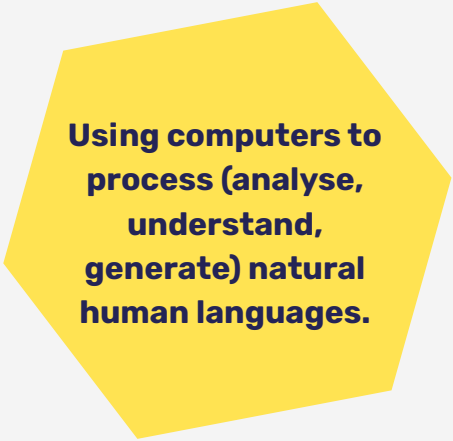

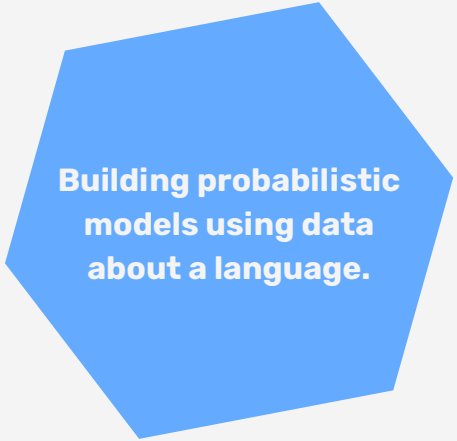

Examples

High Level

Chatbots:
Understand natural language from the user and return intelligent responses.

➔ Api.ai

Information retrieval:
Find relevant results and similar results.

➔ Google

Information extraction:
Structured information from unstructured documents.

➔ Events from Gmail

Machine translation:
One language to another.

➔ Google Translate

Text simplification:
Preserve the meaning of text, but simplify the grammar and vocabulary.

➔ Rewordify

➔ Simple English Wikipedia

Predictive text input:
Faster or easier typing.

➔ Phrase completion application

➔ A much better application

Sentiment analysis:
Attitude of speaker.

➔ Hater News

Speech recognition and generation:
Speech-to-text, text-to-speech.

➔ Google's Web Speech API demo

➔ Vocalware Text-to-Speech demo

Low Level

Tokenization:
Breaking text into tokens (words, sentences, n-grams)

Stop-word removal:
a/an/the

Stemming and lemmatization:
root word

TF-IDF:
word importance

Part-of-speech tagging:
noun/verb/adjective

Named entity recognition:
person/organization /location

Spelling correction:
"New Yrok City"

Word sense disambiguation:
"buy a mouse"

Segmentation:
"New York City subway"

Language detection:
"translate this page"

Machine learning:
specialized models that work well with text

Challenges

Ambiguity:

- ➔ Hospitals Are Sued by 7 Foot Doctors
- ➔ Juvenile Court to Try Shooting Defendant
- ➔ Local High School Dropouts Cut in Half

Non-standard English:
text messages

Idioms:
"throw in the towel"

Newly coined words:
"retweet"

Tricky entity names:
"Where is A Bug's Life playing?"

World knowledge:
"Mary and Sue are sisters", "Mary and Sue are mothers"

multiverse

Text Classification

Text Classification

Text Classification

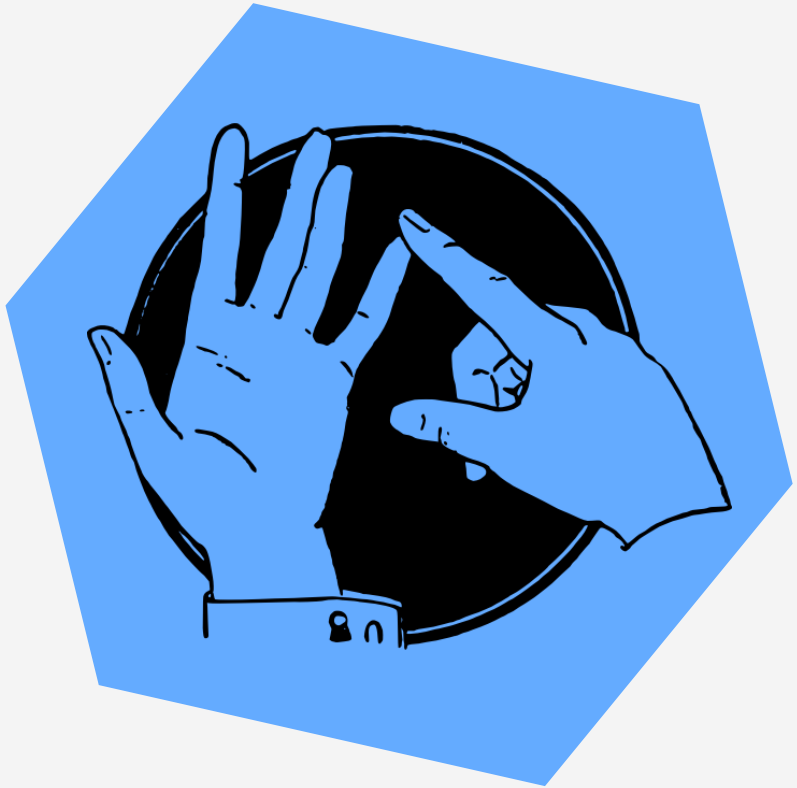

Text Classification

Text Classification

Text Classification

```
from sklearn.feature_extraction.text import CountVectorizer
# Use CountVectorizer to create document-term matrices from X_train and X_test.
vect = CountVectorizer()
X_train_dtm = vect.fit_transform(X_train)
X_test_dtm = vect.transform(X_test)
```

Text Classification

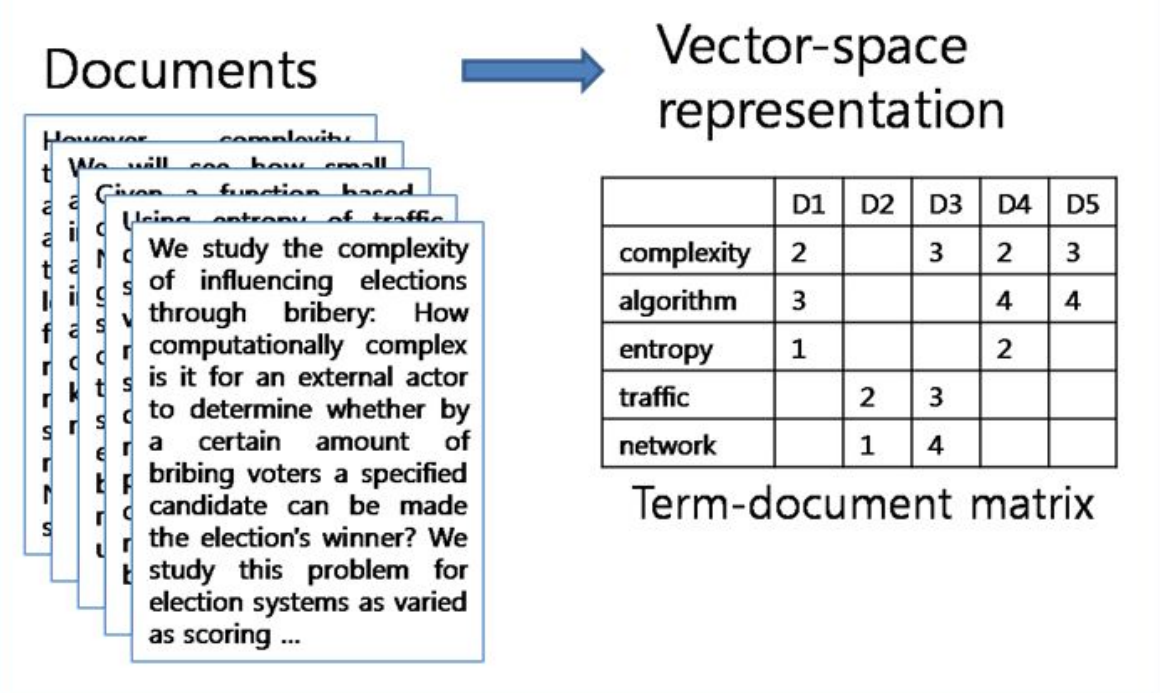

	D1	D2	D3	D4	D5
complexity	2		3	2	3
algorithm	3			4	4
entropy	1			2	
traffic		2	3		
network		1	4		

multiverse

Let’s Practice

multiverse

NGrams

NGrams

```
my cat is awesome
Unigrams (1-grams): 'my', 'cat', 'is', 'awesome'
Bigrams (2-grams): 'my cat', 'cat is', 'is awesome'
Trigrams (3-grams): 'my cat is', 'cat is awesome'
4-grams: 'my cat is awesome'
```

NGrams

```
# Include 1-grams and 2-grams.
vect = CountVectorizer(ngram_range=(1, 2))
X_train_dtm = vect.fit_transform(X_train)
X_train_dtm.shape
```

multiverse

Let’s Practice

multiverse

Stop Words and Other Options

Stop Words

Stop Words

```
# Remove English stop words.
vect = CountVectorizer(stop_words='english')
```

MAX

```
vect = CountVectorizer(max_df=0.7)
```

```
vect = CountVectorizer(max_features=100)
```

multiverse

Let’s Practice

multiverse

Stemming and Lemmatization

Stemming

Lemmatization

Lemmatization	Stemming
shouted – shout	badly – bad
best – good	computing – comput
better – good	computed – comput
good – good	wipes – wip
wiping – wipe	wiped – wip
hidden – hide	wiping – wip

multiverse

Let’s Practice

multiverse

TF-IDF

TF-IDF

```
# TfidfVectorizer
vect = TfidfVectorizer()
pd.DataFrame(vect.fit_transform(simple_train).toarray(), columns=vect.get_feature_names())
```

	cab	call	me	please	tonight	you
0	0.000000	0.385372	0.000000	0.000000	0.652491	0.652491
1	0.720333	0.425441	0.547832	0.000000	0.000000	0.000000
2	0.000000	0.266075	0.342620	0.901008	0.000000	0.000000

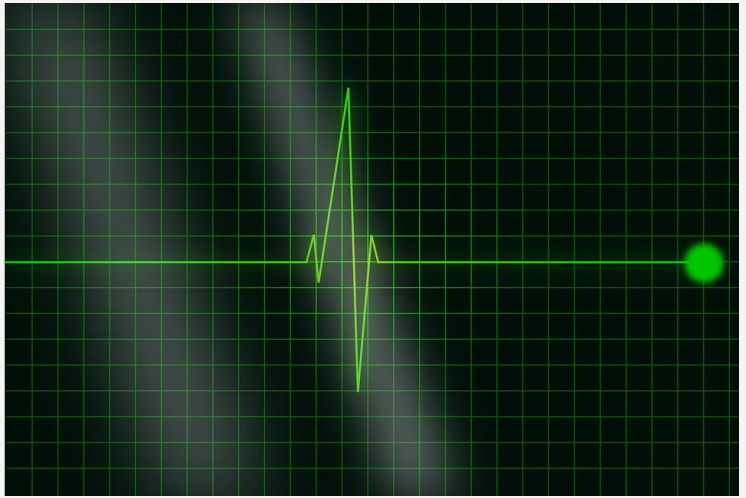

multiverse

Let’s Practice

multiverse

Sentiment Analysis

I Appreciate the Sentiment

I Appreciate the Sentiment

multiverse

Let’s Practice